

O-MAPL: Offline Multi-Agent Preference Learning

This repository contains the implementation of **O-MAPL** (Offline Multi-Agent Preference Learning), an end-to-end framework for preference-based learning in cooperative multi-agent reinforcement learning (MARL). O-MAPL eliminates the need for explicit reward modeling by directly learning Q-functions from human preferences. The framework is designed to handle both rule-based and LLM-based preference datasets, with experiments conducted on **SMACv1**, **SMACv2**, and **MaMujoco** benchmarks.

Table of Contents

- [Overview](#)
- [Project Structure](#)
- [Installation Instructions](#)
- [Usage](#)
 - [Step 1: Generate LLM Dataset](#)
 - [Step 2: Training](#)
 - [Step 3: Evaluation](#)
 - [Step 4: Analysis and Visualization](#)
- [Datasets](#)
 - [Rule-Based Preference Data](#)
 - [LLM-Based Preference Data](#)
- [Key Features](#)
- [Results](#)
- [License](#)

Overview

O-MAPL introduces a unified, end-to-end learning process for multi-agent preference-based reinforcement learning. Instead of relying on a two-phase approach (reward modeling followed by policy optimization), O-MAPL directly learns Q-functions from preference data. Key highlights include:

- **Value Factorization:** A novel value decomposition strategy ensures scalability and stability in multi-agent systems.
- **LLM Integration:** Support for preference datasets generated using large language models (LLMs) like GPT-4.
- **Benchmarks:** Comprehensive evaluation on **SMACv1**, **SMACv2**, and **MaMujoco** benchmarks.

Project Structure

```
.
├── algos                               # Implementation of O-MAPL and baseline algorithms
│   ├── bc.py                          # Behavioral Cloning (BC)
│   ├── iipl.py                        # Independent Inverse Preference Learning (IIPL)
│   ├── iplvdn.py                      # IPL with VDN aggregation
│   ├── omarl.py                       # O-MAPL implementation
│   ├── slmarl.py                      # Supervised Learning for MARL (2-phase approach)
│   └── utils.py                       # Utilities for algorithms
├── envs                                # Environment wrappers for SMACv1, SMACv2, and MaMujoco
│   ├── mamujoco                       # Multi-agent MuJoCo environments
│   ├── smacv1                         # SMACv1 environments
│   └── smacv2                         # SMACv2 environments
├── trainers                           # Training scripts for each algorithm
├── rollouts                           # Rollout utilities for continuous and discrete environments
├── saved_results_final                 # Saved experimental results (rule-based and LLM-based)
├── graphs                             # Pre-generated graphs for analysis
├── analyze.ipynb                      # Jupyter notebook for analysis and visualization
├── llm_generate.py                    # Script to generate prompts for LLM-based annotations
├── llm_upload.py                      # Script to upload prompts to OpenAI APIs
├── llm_retrieve.py                    # Script to retrieve LLM-based annotation results
└── main.py                            # Main training script
```

└─ evaluate.py	# Evaluation script
└─ config.py	# Configuration file
└─ utils.py	# General utilities
└─ readme.md	# Readme file

Installation Instructions

Follow these steps to set up and run the project:

1. Clone the repository:

```
git clone https://github.com/your-repo/omapl.git
cd omapl
```

2. Install dependencies: Ensure you have Python 3.8+ installed. Then, install the required packages:

```
pip install -r requirements.txt
```

3. Set up OpenAI API access: For generating LLM-based datasets, ensure you have access to OpenAI APIs. Set your API key as an environment variable:

```
export OPENAI_API_KEY="your-api-key"
```

Usage

Step 1: Generate LLM Dataset

To generate preference datasets using LLM-based annotations, follow these steps:

1. Generate prompts: This script generates prompts based on trajectory data:

```
python -u llm_generate.py
```

2. Upload prompts to OpenAI: Use the llm_upload.py script to submit prompts to the OpenAI API in batches:

```
python -u llm_upload.py
```

This script will handle token usage and batch processing.

3. Retrieve LLM responses: Once the prompts are processed, retrieve the results:

```
python -u llm_retrieve.py
```

The generated preference data will be saved for use in training.

Step 2: Training

Train the O-MAPL algorithm or baselines on different environments:

- **For MaMujoco:**

```
python -u main.py --algo OMAPL --env_name Ant-v2 --seed 0 --lr 1e-4 --action_scale 0.7
```

- **For SMACv1:**

```
python -u main.py --algo OMAPL --env_name 5m_vs_6m --seed 0 --batch_size 8 --lr 1e-4
```

- **For SMACv2:**

```
python -u main.py --algo OMAPL --env_name protoss_5_vs_5 --seed 0 --batch_size 8 --lr 1e-4
```

To train with LLM-based datasets, add the `--use-llm` flag:

```
python -u main.py --algo OMAPL --env_name protoss_5_vs_5 --use-llm
```

Step 3: Evaluation

Evaluate trained models using the `evaluate.py` script:

```
python -u evaluate.py --algo {algo} --env_name {env_name} --seed {seed} --eval_step {eval_step}
```

Replace `{algo}`, `{env_name}`, `{seed}`, and `{eval_step}` with the appropriate algorithm, environment, seed number, and training step.

Step 4: Analysis and Visualization

Use the `analyze.ipynb` notebook for detailed analysis:

- Load experimental results.
- Visualize performance metrics (returns, win rates).
- Generate tables and graphs for evaluation.

Example:

- Rule-based results: `saved_results_final/results.pkl`
- LLM-based results: `saved_results_final/results_llm.pkl`

Datasets

Rule-Based Preference Data

- Generated by sampling trajectories from offline datasets and assigning binary preferences based on quality (e.g., poor vs. expert).
- Covers **SMACv1**, **SMACv2**, and **MaMujoco** environments.

LLM-Based Preference Data

- Generated using GPT-4 with detailed prompts describing trajectory outcomes.
- Provides richer annotations for tasks like **SMACv1** and **SMACv2**.
- Example prompt:

```
Scenario: 5m_vs_6m
Allied Agents Health: [0.5, 0.4, 0.3, 0.2, 0.1]
Enemy Agents Health: [0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
Total Steps: 28
Which trajectory is better? #1 or #2
```

Key Features

- **Algorithms:**
 - O-MAPL (proposed method)
 - Baselines: BC, IIPL, IPL-VDN, SL-MARL
- **Environments:**
 - SMACv1: Cooperative tasks in StarCraft II.
 - SMACv2: Enhanced StarCraft II tasks with randomized start positions and unit types.
 - MaMujoco: Continuous control tasks for multi-agent robotics.
- **Datasets:**
 - Rule-based and LLM-generated preference datasets.
- **Visualization:**
 - Pre-generated graphs for returns and win rates (see `graphs/` folder).

Results

- **SMACv1**: O-MAPL achieves the highest win rates in tasks like 2c_vs_64zg and corridor.
- **SMACv2**: Significant improvements in complex scenarios like protoss_10_vs_11 and terran_20_vs_20.
- **MaMujoco**: O-MAPL outperforms baselines in continuous control tasks like Ant-v2 and HalfCheetah-v2.

For detailed results, see the `saved_results_final` folder and the `graphs` directory.

License

This project is licensed under the MIT License. See the `LICENSE` file for details.